

# Keras -- MLPs on MNIST

References: Applied Ai course Google Stackoverflow

```
In [0]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

```
In [0]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [0]: %matplotlib inline
```

```
In [0]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [39]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [41]: # after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of  
shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of  
shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

```
In [42]: # An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
```

```
In [0]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms Lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
In [44]: # example data point after normlizing  
print(X_train[0])
```

[illegible]

0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686
0.99215686	0.99215686	0.83137255	0.52941176	0.51764706	0.0627451

[illegible]

```
In [45]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0,
0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector :", Y_train[0])
```

Class label of first image : 5  
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

## Softmax classifier

```

In [0]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(W.T . X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

```



```
from keras.models import Sequential
from keras.layers import Dense, Activation
```

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

## Assignment

BN + Relu + Dropout2 Layers

```
In [0]:
```

```
In [48]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormal
         # ization-function-in-keras
         # Multilayer perceptron

         # https://intoli.com/blog/neural-network-initialization/
         # If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condi
         # tion with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
         # h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
         # h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
         # h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 

from keras.layers import Dropout
# import BatchNormalization
from keras.layers.normalization import BatchNormalization

model_assign = Sequential()

model_assign.add(Dense(332, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.042, seed=None)))
model_assign.add(BatchNormalization())
model_assign.add(Dropout(0.5))

model_assign.add(Dense(56, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.071, seed=None)))
model_assign.add(BatchNormalization())
model_assign.add(Dropout(0.5))

model_assign.add(Dense(output_dim, activation='softmax'))

model_assign.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_28 (Dense)	(None, 332)	260620
batch_normalization_11 (Batch Normalization)	(None, 332)	1328
dropout_11 (Dropout)	(None, 332)	0
dense_29 (Dense)	(None, 56)	18648
batch_normalization_12 (Batch Normalization)	(None, 56)	224
dropout_12 (Dropout)	(None, 56)	0
dense_30 (Dense)	(None, 10)	570
=====		
Total params: 281,390		
Trainable params: 280,614		
Non-trainable params: 776		

```
In [49]: model_assign.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_assign.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 69us/step - loss: 0.5452 -  
acc: 0.8381 - val\_loss: 0.1653 - val\_acc: 0.9490

Epoch 2/20

60000/60000 [=====] - 3s 55us/step - loss: 0.2592 -  
acc: 0.9250 - val\_loss: 0.1235 - val\_acc: 0.9626

Epoch 3/20

60000/60000 [=====] - 3s 58us/step - loss: 0.2047 -  
acc: 0.9413 - val\_loss: 0.1064 - val\_acc: 0.9673

Epoch 4/20

60000/60000 [=====] - 4s 59us/step - loss: 0.1713 -  
acc: 0.9503 - val\_loss: 0.0921 - val\_acc: 0.9712

Epoch 5/20

60000/60000 [=====] - 3s 52us/step - loss: 0.1539 -  
acc: 0.9558 - val\_loss: 0.0865 - val\_acc: 0.9729

Epoch 6/20

60000/60000 [=====] - 3s 49us/step - loss: 0.1410 -  
acc: 0.9595 - val\_loss: 0.0819 - val\_acc: 0.9758

Epoch 7/20

60000/60000 [=====] - 3s 49us/step - loss: 0.1304 -  
acc: 0.9616 - val\_loss: 0.0796 - val\_acc: 0.9758

Epoch 8/20

60000/60000 [=====] - 3s 49us/step - loss: 0.1250 -  
acc: 0.9634 - val\_loss: 0.0748 - val\_acc: 0.9765

Epoch 9/20

60000/60000 [=====] - 3s 49us/step - loss: 0.1141 -  
acc: 0.9671 - val\_loss: 0.0684 - val\_acc: 0.9795

Epoch 10/20

60000/60000 [=====] - 3s 49us/step - loss: 0.1057 -  
acc: 0.9690 - val\_loss: 0.0688 - val\_acc: 0.9785

Epoch 11/20

60000/60000 [=====] - 3s 49us/step - loss: 0.1035 -  
acc: 0.9690 - val\_loss: 0.0734 - val\_acc: 0.9782

Epoch 12/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0963 -  
acc: 0.9716 - val\_loss: 0.0689 - val\_acc: 0.9797

Epoch 13/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0931 -  
acc: 0.9721 - val\_loss: 0.0657 - val\_acc: 0.9808

Epoch 14/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0923 -  
acc: 0.9722 - val\_loss: 0.0647 - val\_acc: 0.9798

Epoch 15/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0865 -  
acc: 0.9735 - val\_loss: 0.0701 - val\_acc: 0.9787

Epoch 16/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0876 -  
acc: 0.9734 - val\_loss: 0.0689 - val\_acc: 0.9802

Epoch 17/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0810 -  
acc: 0.9762 - val\_loss: 0.0666 - val\_acc: 0.9803

Epoch 18/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0786 -  
acc: 0.9764 - val\_loss: 0.0690 - val\_acc: 0.9808

Epoch 19/20

60000/60000 [=====] - 3s 49us/step - loss: 0.0788 -

acc: 0.9763 - val\_loss: 0.0681 - val\_acc: 0.9798  
 Epoch 20/20  
 60000/60000 [=====] - 3s 49us/step - loss: 0.0744 -  
 acc: 0.9776 - val\_loss: 0.0661 - val\_acc: 0.9812

```
In [50]: score = model_assign.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06606815920227091

Test accuracy: 0.9812



### 3 Hidden layers

```

In [51]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormal
         # ization-function-in-keras
         # Multilayer perceptron

         # https://intoli.com/blog/neural-network-initialization/
         # If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condi
         # tion with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
         # h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
         # h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
         # h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers import Dropout

model_assign3 = Sequential()

model_assign3.add(Dense(512, activation='relu', input_shape=(input_dim,), kern
el_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_assign3.add(BatchNormalization())
model_assign3.add(Dropout(0.5))

model_assign3.add(Dense(128, activation='relu', input_shape=(input_dim,), kern
el_initializer=RandomNormal(mean=0.0, stddev=0.056, seed=None)))
model_assign3.add(BatchNormalization())
model_assign3.add(Dropout(0.5))

model_assign3.add(Dense(64, activation='relu', kernel_initializer=RandomNormal
(mean=0.0, stddev=0.10, seed=None)) )
model_assign3.add(BatchNormalization())
model_assign3.add(Dropout(0.5))

model_assign3.add(Dense(output_dim, activation='softmax'))

model_assign3.summary()

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_31 (Dense)	(None, 512)	401920
batch_normalization_13 (Batch Normalization)	(None, 512)	2048
dropout_13 (Dropout)	(None, 512)	0
dense_32 (Dense)	(None, 128)	65664
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dropout_14 (Dropout)	(None, 128)	0
dense_33 (Dense)	(None, 64)	8256
batch_normalization_15 (Batch Normalization)	(None, 64)	256
dropout_15 (Dropout)	(None, 64)	0
dense_34 (Dense)	(None, 10)	650
=====	=====	=====
Total params: 479,306		
Trainable params: 477,898		
Non-trainable params: 1,408		
=====	=====	=====

```
In [52]: model_assign3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_assign3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 88us/step - loss: 0.6582 -  
acc: 0.8005 - val\_loss: 0.1684 - val\_acc: 0.9460

Epoch 2/20

60000/60000 [=====] - 4s 70us/step - loss: 0.2803 -  
acc: 0.9203 - val\_loss: 0.1247 - val\_acc: 0.9613

Epoch 3/20

60000/60000 [=====] - 4s 71us/step - loss: 0.2149 -  
acc: 0.9399 - val\_loss: 0.1083 - val\_acc: 0.9669

Epoch 4/20

60000/60000 [=====] - 4s 66us/step - loss: 0.1845 -  
acc: 0.9487 - val\_loss: 0.0929 - val\_acc: 0.9711

Epoch 5/20

60000/60000 [=====] - 4s 61us/step - loss: 0.1624 -  
acc: 0.9553 - val\_loss: 0.0888 - val\_acc: 0.9718

Epoch 6/20

60000/60000 [=====] - 4s 61us/step - loss: 0.1463 -  
acc: 0.9585 - val\_loss: 0.0820 - val\_acc: 0.9760

Epoch 7/20

60000/60000 [=====] - 4s 61us/step - loss: 0.1341 -  
acc: 0.9629 - val\_loss: 0.0838 - val\_acc: 0.9747

Epoch 8/20

60000/60000 [=====] - 4s 61us/step - loss: 0.1268 -  
acc: 0.9642 - val\_loss: 0.0781 - val\_acc: 0.9755

Epoch 9/20

60000/60000 [=====] - 4s 62us/step - loss: 0.1165 -  
acc: 0.9673 - val\_loss: 0.0716 - val\_acc: 0.9783

Epoch 10/20

60000/60000 [=====] - 4s 62us/step - loss: 0.1107 -  
acc: 0.9688 - val\_loss: 0.0700 - val\_acc: 0.9801

Epoch 11/20

60000/60000 [=====] - 4s 61us/step - loss: 0.1043 -  
acc: 0.9700 - val\_loss: 0.0725 - val\_acc: 0.9788

Epoch 12/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0998 -  
acc: 0.9714 - val\_loss: 0.0661 - val\_acc: 0.9808

Epoch 13/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0959 -  
acc: 0.9729 - val\_loss: 0.0634 - val\_acc: 0.9815

Epoch 14/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0918 -  
acc: 0.9741 - val\_loss: 0.0617 - val\_acc: 0.9819

Epoch 15/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0864 -  
acc: 0.9750 - val\_loss: 0.0657 - val\_acc: 0.9819

Epoch 16/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0845 -  
acc: 0.9763 - val\_loss: 0.0633 - val\_acc: 0.9822

Epoch 17/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0836 -  
acc: 0.9757 - val\_loss: 0.0632 - val\_acc: 0.9825

Epoch 18/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0791 -  
acc: 0.9767 - val\_loss: 0.0659 - val\_acc: 0.9826

Epoch 19/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0792 -

acc: 0.9778 - val\_loss: 0.0614 - val\_acc: 0.9837  
 Epoch 20/20  
 60000/60000 [=====] - 4s 61us/step - loss: 0.0755 -  
 acc: 0.9780 - val\_loss: 0.0672 - val\_acc: 0.9813

```
In [53]: score = model_assign3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06720245788121829

Test accuracy: 0.9813



5 hidden layers

```

In [54]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormal
         # ization-function-in-keras
         # Multilayer perceptron

         # https://intoli.com/blog/neural-network-initialization/
         # If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condi
         # tion with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
         # h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
         # h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
         # h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers import Dropout

model_assign5 = Sequential()

model_assign5.add(Dense(512, activation='relu', input_shape=(input_dim,), kern
el_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_assign5.add(BatchNormalization())
model_assign5.add(Dropout(0.5))

model_assign5.add(Dense(256, activation='relu', kernel_initializer=RandomNormal
(mean=0.0, stddev=0.051, seed=None)) )
model_assign5.add(BatchNormalization())
model_assign5.add(Dropout(0.5))

model_assign5.add(Dense(128, activation='relu', kernel_initializer=RandomNormal
(mean=0.0, stddev=0.072, seed=None)) )
model_assign5.add(BatchNormalization())
model_assign5.add(Dropout(0.5))

model_assign5.add(Dense(64, activation='relu', kernel_initializer=RandomNormal
(mean=0.0, stddev=0.10, seed=None)) )
model_assign5.add(BatchNormalization())
model_assign5.add(Dropout(0.5))

model_assign5.add(Dense(32, activation='relu', kernel_initializer=RandomNormal
(mean=0.0, stddev=0.144, seed=None)) )
model_assign5.add(BatchNormalization())
model_assign5.add(Dropout(0.5))

model_assign5.add(Dense(output_dim, activation='softmax'))

model_assign5.summary()

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_35 (Dense)	(None, 512)	401920
batch_normalization_16 (Batch Normalization)	(None, 512)	2048
dropout_16 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 256)	131328
batch_normalization_17 (Batch Normalization)	(None, 256)	1024
dropout_17 (Dropout)	(None, 256)	0
dense_37 (Dense)	(None, 128)	32896
batch_normalization_18 (Batch Normalization)	(None, 128)	512
dropout_18 (Dropout)	(None, 128)	0
dense_38 (Dense)	(None, 64)	8256
batch_normalization_19 (Batch Normalization)	(None, 64)	256
dropout_19 (Dropout)	(None, 64)	0
dense_39 (Dense)	(None, 32)	2080
batch_normalization_20 (Batch Normalization)	(None, 32)	128
dropout_20 (Dropout)	(None, 32)	0
dense_40 (Dense)	(None, 10)	330
=====	=====	=====
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		
=====	=====	=====

```
In [55]: model_assign5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_assign5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 124us/step - loss: 1.2889 - acc: 0.5871 - val\_loss: 0.2800 - val\_acc: 0.9286

Epoch 2/20

60000/60000 [=====] - 6s 99us/step - loss: 0.5212 - acc: 0.8593 - val\_loss: 0.1755 - val\_acc: 0.9521

Epoch 3/20

60000/60000 [=====] - 6s 93us/step - loss: 0.3722 - acc: 0.9079 - val\_loss: 0.1519 - val\_acc: 0.9609

Epoch 4/20

60000/60000 [=====] - 5s 86us/step - loss: 0.3107 - acc: 0.9253 - val\_loss: 0.1302 - val\_acc: 0.9677

Epoch 5/20

60000/60000 [=====] - 5s 86us/step - loss: 0.2724 - acc: 0.9364 - val\_loss: 0.1224 - val\_acc: 0.9663

Epoch 6/20

60000/60000 [=====] - 5s 86us/step - loss: 0.2483 - acc: 0.9423 - val\_loss: 0.1099 - val\_acc: 0.9738

Epoch 7/20

60000/60000 [=====] - 5s 86us/step - loss: 0.2243 - acc: 0.9488 - val\_loss: 0.1181 - val\_acc: 0.9716

Epoch 8/20

60000/60000 [=====] - 5s 88us/step - loss: 0.2091 - acc: 0.9518 - val\_loss: 0.1143 - val\_acc: 0.9731

Epoch 9/20

60000/60000 [=====] - 6s 96us/step - loss: 0.1986 - acc: 0.9546 - val\_loss: 0.1015 - val\_acc: 0.9741

Epoch 10/20

60000/60000 [=====] - 5s 85us/step - loss: 0.1906 - acc: 0.9569 - val\_loss: 0.0883 - val\_acc: 0.9782

Epoch 11/20

60000/60000 [=====] - 5s 85us/step - loss: 0.1798 - acc: 0.9594 - val\_loss: 0.0958 - val\_acc: 0.9770

Epoch 12/20

60000/60000 [=====] - 5s 85us/step - loss: 0.1719 - acc: 0.9615 - val\_loss: 0.0883 - val\_acc: 0.9786

Epoch 13/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1658 - acc: 0.9629 - val\_loss: 0.0887 - val\_acc: 0.9792

Epoch 14/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1586 - acc: 0.9647 - val\_loss: 0.0900 - val\_acc: 0.9789

Epoch 15/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1491 - acc: 0.9662 - val\_loss: 0.0851 - val\_acc: 0.9799

Epoch 16/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1410 - acc: 0.9681 - val\_loss: 0.0894 - val\_acc: 0.9796

Epoch 17/20

60000/60000 [=====] - 6s 99us/step - loss: 0.1473 - acc: 0.9672 - val\_loss: 0.0963 - val\_acc: 0.9793

Epoch 18/20

60000/60000 [=====] - 6s 96us/step - loss: 0.1407 - acc: 0.9677 - val\_loss: 0.0853 - val\_acc: 0.9798

Epoch 19/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1344 -

acc: 0.9695 - val\_loss: 0.0755 - val\_acc: 0.9823  
 Epoch 20/20  
 60000/60000 [=====] - 5s 86us/step - loss: 0.1291 -  
 acc: 0.9703 - val\_loss: 0.0864 - val\_acc: 0.9796

```
In [56]: score = model_assign5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08644420035893563

Test accuracy: 0.9796



# Without Dropout and BN

2 layers

```
In [57]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormal
         # ization-function-in-keras
         # Multilayer perceptron

         # https://intoli.com/blog/neural-network-initialization/
         # If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condi
         # tion with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
         # h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
         # h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
         # h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

         from keras.layers import Dropout

         model_assign = Sequential()

         model_assign.add(Dense(332, activation='relu', input_shape=(input_dim,), kernel_
         l_initializer=RandomNormal(mean=0.0, stddev=0.042, seed=None)))

         model_assign.add(Dense(56, activation='relu', kernel_initializer=RandomNormal(
         mean=0.0, stddev=0.071, seed=None)) )

         model_assign.add(Dense(output_dim, activation='softmax'))

         model_assign.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_41 (Dense)	(None, 332)	260620
dense_42 (Dense)	(None, 56)	18648
dense_43 (Dense)	(None, 10)	570
=====		
Total params: 279,838		
Trainable params: 279,838		
Non-trainable params: 0		
=====		



```
In [58]: model_assign.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_assign.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 55us/step - loss: 0.2737 -  
acc: 0.9219 - val\_loss: 0.1348 - val\_acc: 0.9589

Epoch 2/20

60000/60000 [=====] - 2s 31us/step - loss: 0.1079 -  
acc: 0.9680 - val\_loss: 0.0925 - val\_acc: 0.9705

Epoch 3/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0686 -  
acc: 0.9786 - val\_loss: 0.0763 - val\_acc: 0.9774

Epoch 4/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0489 -  
acc: 0.9850 - val\_loss: 0.0731 - val\_acc: 0.9782

Epoch 5/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0376 -  
acc: 0.9880 - val\_loss: 0.0758 - val\_acc: 0.9771

Epoch 6/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0268 -  
acc: 0.9921 - val\_loss: 0.0680 - val\_acc: 0.9793

Epoch 7/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0215 -  
acc: 0.9934 - val\_loss: 0.0743 - val\_acc: 0.9778

Epoch 8/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0166 -  
acc: 0.9950 - val\_loss: 0.0738 - val\_acc: 0.9789

Epoch 9/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0151 -  
acc: 0.9951 - val\_loss: 0.0911 - val\_acc: 0.9767

Epoch 10/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0121 -  
acc: 0.9960 - val\_loss: 0.0688 - val\_acc: 0.9812

Epoch 11/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0123 -  
acc: 0.9961 - val\_loss: 0.0843 - val\_acc: 0.9808

Epoch 12/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0094 -  
acc: 0.9969 - val\_loss: 0.0737 - val\_acc: 0.9824

Epoch 13/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0104 -  
acc: 0.9966 - val\_loss: 0.0800 - val\_acc: 0.9809

Epoch 14/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0061 -  
acc: 0.9981 - val\_loss: 0.0852 - val\_acc: 0.9790

Epoch 15/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0110 -  
acc: 0.9961 - val\_loss: 0.0917 - val\_acc: 0.9781

Epoch 16/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0086 -  
acc: 0.9971 - val\_loss: 0.0967 - val\_acc: 0.9776

Epoch 17/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0091 -  
acc: 0.9970 - val\_loss: 0.0804 - val\_acc: 0.9811

Epoch 18/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0069 -  
acc: 0.9974 - val\_loss: 0.0824 - val\_acc: 0.9802

Epoch 19/20

60000/60000 [=====] - 2s 31us/step - loss: 0.0025 -

acc: 0.9993 - val\_loss: 0.0865 - val\_acc: 0.9822  
 Epoch 20/20  
 60000/60000 [=====] - 2s 31us/step - loss: 0.0028 -  
 acc: 0.9991 - val\_loss: 0.1038 - val\_acc: 0.9789

```
In [59]: score = model_assign.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1037857443765759

Test accuracy: 0.9789



3 Layers

```
In [60]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormal
         # ization-function-in-keras
         # Multilayer perceptron

         # https://intoli.com/blog/neural-network-initialization/
         # If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condi
         # tion with  $\sigma=\sqrt{2/(n_i+n_i+1)}$ .
         # h1 =>  $\sigma=\sqrt{2/(n_i+n_i+1)} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
         # h2 =>  $\sigma=\sqrt{2/(n_i+n_i+1)} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
         # h1 =>  $\sigma=\sqrt{2/(n_i+n_i+1)} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers import Dropout

model_assign3 = Sequential()

model_assign3.add(Dense(512, activation='relu', input_shape=(input_dim,), kern
el_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))

model_assign3.add(Dense(128, activation='relu', input_shape=(input_dim,), kern
el_initializer=RandomNormal(mean=0.0, stddev=0.056, seed=None)))

model_assign3.add(Dense(64, activation='relu', kernel_initializer=RandomNormal
(mean=0.0, stddev=0.10, seed=None)) )

model_assign3.add(Dense(output_dim, activation='softmax'))

model_assign3.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_44 (Dense)	(None, 512)	401920
dense_45 (Dense)	(None, 128)	65664
dense_46 (Dense)	(None, 64)	8256
dense_47 (Dense)	(None, 10)	650
=====	=====	=====
Total params: 476,490		
Trainable params: 476,490		
Non-trainable params: 0		
=====	=====	=====

```
In [61]: model_assign3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_assign3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 61us/step - loss: 0.2583 -  
acc: 0.9237 - val\_loss: 0.1316 - val\_acc: 0.9590

Epoch 2/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0916 -  
acc: 0.9723 - val\_loss: 0.0840 - val\_acc: 0.9730

Epoch 3/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0591 -  
acc: 0.9815 - val\_loss: 0.0760 - val\_acc: 0.9766

Epoch 4/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0409 -  
acc: 0.9874 - val\_loss: 0.0777 - val\_acc: 0.9755

Epoch 5/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0329 -  
acc: 0.9889 - val\_loss: 0.0730 - val\_acc: 0.9784

Epoch 6/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0249 -  
acc: 0.9918 - val\_loss: 0.0742 - val\_acc: 0.9788

Epoch 7/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0208 -  
acc: 0.9930 - val\_loss: 0.0672 - val\_acc: 0.9799

Epoch 8/20

60000/60000 [=====] - 2s 35us/step - loss: 0.0164 -  
acc: 0.9945 - val\_loss: 0.0803 - val\_acc: 0.9793

Epoch 9/20

60000/60000 [=====] - 2s 39us/step - loss: 0.0169 -  
acc: 0.9944 - val\_loss: 0.0854 - val\_acc: 0.9787

Epoch 10/20

60000/60000 [=====] - 2s 39us/step - loss: 0.0185 -  
acc: 0.9939 - val\_loss: 0.0790 - val\_acc: 0.9795

Epoch 11/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0117 -  
acc: 0.9960 - val\_loss: 0.0717 - val\_acc: 0.9828

Epoch 12/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0110 -  
acc: 0.9964 - val\_loss: 0.0912 - val\_acc: 0.9795

Epoch 13/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0104 -  
acc: 0.9965 - val\_loss: 0.0808 - val\_acc: 0.9820

Epoch 14/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0096 -  
acc: 0.9968 - val\_loss: 0.0994 - val\_acc: 0.9770

Epoch 15/20

60000/60000 [=====] - 2s 35us/step - loss: 0.0112 -  
acc: 0.9962 - val\_loss: 0.0861 - val\_acc: 0.9808

Epoch 16/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0131 -  
acc: 0.9956 - val\_loss: 0.0843 - val\_acc: 0.9802

Epoch 17/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0081 -  
acc: 0.9976 - val\_loss: 0.0912 - val\_acc: 0.9802

Epoch 18/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0077 -  
acc: 0.9974 - val\_loss: 0.0811 - val\_acc: 0.9831

Epoch 19/20

60000/60000 [=====] - 2s 34us/step - loss: 0.0084 -

acc: 0.9972 - val\_loss: 0.0767 - val\_acc: 0.9833  
 Epoch 20/20  
 60000/60000 [=====] - 2s 34us/step - loss: 0.0068 -  
 acc: 0.9979 - val\_loss: 0.0917 - val\_acc: 0.9829

```
In [62]: score = model_assign3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09172725084941308

Test accuracy: 0.9829



```
In [63]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormal
         # ization-function-in-keras
         # Multilayer perceptron

         # https://intoli.com/blog/neural-network-initialization/
         # If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condi
         # tion with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
         # h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
         # h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
         # h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers import Dropout

model_assign5 = Sequential()

model_assign5.add(Dense(512, activation='relu', input_shape=(input_dim,), kern
el_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))

model_assign5.add(Dense(256, activation='relu', kernel_initializer=RandomNorma
l(mean=0.0, stddev=0.051, seed=None)) )

model_assign5.add(Dense(128, activation='relu', kernel_initializer=RandomNorma
l(mean=0.0, stddev=0.072, seed=None)) )

model_assign5.add(Dense(64, activation='relu', kernel_initializer=RandomNormal
(mean=0.0, stddev=0.10, seed=None)) )

model_assign5.add(Dense(32, activation='relu', kernel_initializer=RandomNormal
(mean=0.0, stddev=0.144, seed=None)) )

model_assign5.add(Dense(output_dim, activation='softmax'))

model_assign5.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_48 (Dense)	(None, 512)	401920
dense_49 (Dense)	(None, 256)	131328
dense_50 (Dense)	(None, 128)	32896
dense_51 (Dense)	(None, 64)	8256
dense_52 (Dense)	(None, 32)	2080
dense_53 (Dense)	(None, 10)	330
=====	=====	=====
Total params: 576,810		
Trainable params: 576,810		
Non-trainable params: 0		
=====		



```
In [64]: model_assign5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_assign5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 67us/step - loss: 0.2572 -  
acc: 0.9217 - val\_loss: 0.1074 - val\_acc: 0.9669

Epoch 2/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0914 -  
acc: 0.9713 - val\_loss: 0.0969 - val\_acc: 0.9664

Epoch 3/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0617 -  
acc: 0.9806 - val\_loss: 0.0712 - val\_acc: 0.9782

Epoch 4/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0471 -  
acc: 0.9852 - val\_loss: 0.0796 - val\_acc: 0.9769

Epoch 5/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0366 -  
acc: 0.9888 - val\_loss: 0.0903 - val\_acc: 0.9745

Epoch 6/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0332 -  
acc: 0.9895 - val\_loss: 0.0828 - val\_acc: 0.9766

Epoch 7/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0274 -  
acc: 0.9916 - val\_loss: 0.0801 - val\_acc: 0.9785

Epoch 8/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0233 -  
acc: 0.9925 - val\_loss: 0.0763 - val\_acc: 0.9806

Epoch 9/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0203 -  
acc: 0.9939 - val\_loss: 0.0910 - val\_acc: 0.9793

Epoch 10/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0190 -  
acc: 0.9940 - val\_loss: 0.0978 - val\_acc: 0.9784

Epoch 11/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0184 -  
acc: 0.9943 - val\_loss: 0.1111 - val\_acc: 0.9741

Epoch 12/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0146 -  
acc: 0.9954 - val\_loss: 0.0886 - val\_acc: 0.9806

Epoch 13/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0163 -  
acc: 0.9947 - val\_loss: 0.0835 - val\_acc: 0.9832

Epoch 14/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0126 -  
acc: 0.9961 - val\_loss: 0.0901 - val\_acc: 0.9806

Epoch 15/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0132 -  
acc: 0.9961 - val\_loss: 0.0800 - val\_acc: 0.9816

Epoch 16/20

60000/60000 [=====] - 2s 39us/step - loss: 0.0127 -  
acc: 0.9963 - val\_loss: 0.0900 - val\_acc: 0.9823

Epoch 17/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0115 -  
acc: 0.9967 - val\_loss: 0.0955 - val\_acc: 0.9808

Epoch 18/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0125 -  
acc: 0.9963 - val\_loss: 0.0914 - val\_acc: 0.9800

Epoch 19/20

60000/60000 [=====] - 2s 38us/step - loss: 0.0088 -

acc: 0.9974 - val\_loss: 0.0976 - val\_acc: 0.9801  
 Epoch 20/20  
 60000/60000 [=====] - 2s 38us/step - loss: 0.0100 -  
 acc: 0.9969 - val\_loss: 0.0905 - val\_acc: 0.9826

```
In [65]: score = model_assign5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_
data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to num
ber of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0904728904638242

Test accuracy: 0.9826



Pretty Table

```
In [2]: # http://zetcode.com/python/prettytable/
from prettytable import PrettyTable

#If you get a ModuleNotFoundError error , install prettytable using: pip3 inst
all prettytable

x = PrettyTable()
x.field_names = ["Model", "Training accuracy", "Test Accuracy", "val_loss"]

x.add_row(["MLP(BN+Relu+Dropout) 2 layer", 0.977, 0.98, 0.0661])
x.add_row(["MLP(BN+Relu+Dropout) 3 layer", 0.978, 0.98, 0.0672])
x.add_row(["MLP(BN+Relu+Dropout) 5 layer", 0.9703, 0.9796, 0.0864])
x.add_row(["MLP(Relu) 2 layer", 0.999, 0.97, 0.103])
x.add_row(["MLP(Relu) 3 layer", 0.998, 0.9829, 0.0917])
x.add_row(["MLP(Relu) 5 layer", 0.997, 0.98, 0.0905])
print(x)
```

```
+-----+-----+-----+-----+
-+
|          Model          | Training accuracy | Test Accuracy | val_loss
|
+-----+-----+-----+-----+
-+
| MLP(BN+Relu+Dropout) 2 layer |      0.977      |      0.98      | 0.0661
|
| MLP(BN+Relu+Dropout) 3 layer |      0.978      |      0.98      | 0.0672
|
| MLP(BN+Relu+Dropout) 5 layer |     0.9703      |     0.9796     | 0.0864
|
|      MLP(Relu) 2 layer      |      0.999      |      0.97      | 0.103
|
|      MLP(Relu) 3 layer      |      0.998      |     0.9829     | 0.0917
|
|      MLP(Relu) 5 layer      |      0.997      |      0.98      | 0.0905
|
+-----+-----+-----+-----+
-+
```

Conclusion From the pretty table we can observe that 1. loss is minimum when model is 2 layer with BN\_Relu+Dropout 2. model has maximum accuracy when mlp with 2 layer 3. when we play with different parameters of mlp like BN Droupout our loss will decrease Steps Taken 1. import the necessary library 2. Import mnist dataset 3. loading train+ test data mnist 4. normalization 5. Giving Output a 10 class 7. Applying models with different parameters